Abbring & Van den Berg (2003), built up step by step

Plan

1. One duration: survival, hazard, integrated hazard
2. People differ: unobserved heterogeneity and dynamic selection
3. The mixed proportional hazard (MPH) model
4. Two durations: competing risks
5. Adding a treatment: potential outcomes indexed by treatment time, no anticipation
6. What we actually observe
7. The model of the paper (Model 1a)
8. Identification, step by step
9. Extensions: heterogeneous effects, multiple spells
10. Connections: DiD with staggered adoption, sequential exogeneity, traps

Running example (the paper's, Sections 1 and 2.4)

- Unemployed workers receive unemployment benefits.
- At a random time S during the spell, their **benefits are reduced** (from level w to $b < w$).
- Outcome: unemployment duration Y.
- Question: what is the causal effect of the benefits reduction on the **transition rate to work**?

Road map: why each section is needed

- The effect is on a *rate* ⇒ we need the language of rates (Section 1).
- Workers differ in ways we do not see, and those still unemployed are selected ⇒ Sections 2–3.
- Before treatment, two things compete: finding a job vs. the benefits reduction ⇒ Section 4.
- Define the causal effect and what we observe ⇒ Sections 5–6.
- Put it together and show the effect is identified ⇒ Sections 7–8.

1 One duration

→ *Why now:* the paper's effect is on a *transition rate*. Before anything else: what is a rate for one duration, and how does it relate to probabilities?

Objects

- $Y \geq 0$: a continuous duration; y: a value.
- Density $f(y)$, CDF $F(y)$.
- Survival function $S(y) := \Pr(Y > y) = 1 - F(y)$: share still in the state at y.

1.1 Hazard rate

→ *Why now:* this is the object the treatment effect will act on.

$$\begin{aligned}
\theta(y) &:= \lim_{dy \downarrow 0} \frac{\Pr(y \le Y < y + dy \mid Y \ge y)}{dy} && \text{[definition]} \\
&= \lim_{dy \downarrow 0} \frac{\Pr(y \le Y < y + dy)}{dy \ \Pr(Y \ge y)} && \text{[conditional probability: } \Pr(A \mid B) = \Pr(A \cap B)/\Pr(B)\text{]} \\
&= \lim_{dy \downarrow 0} \frac{F(y + dy) - F(y)}{dy} \cdot \frac{1}{S(y)} && \text{[}Y \text{ continuous: } \Pr(Y \ge y) = \Pr(Y > y)\text{]} \\
&= \frac{f(y)}{S(y)} && \text{[definition of the derivative: } F' = f\text{]} \qquad (1)
\end{aligned}$$

- Exit rate *among those still in the state* at y.
- $\theta(y)\,dy \approx$ probability of leaving in $[y, y + dy)$, given still there at y.
- A rate, not a probability: can exceed 1.

1.2 Integrated hazard and survival

→ *Why now:* we model hazards, but data and likelihoods are about probabilities. We need the map from θ back to S and f; it is used in every later derivation.

$$\begin{aligned}
\theta(y) &= \frac{f(y)}{S(y)} && \text{[by (1)]} \\
&= -\frac{S'(y)}{S(y)} && \text{[}S = 1 - F\text{, so } S' = -f\text{]} \\
&= -\frac{d}{dy} \log S(y) && \text{[chain rule: } (\log S)' = S'/S\text{]}
\end{aligned}$$

$$\begin{aligned}
\Theta(y) := \int_0^y \theta(u)\,du &= -\big[\log S(u)\big]_0^y && \text{[fundamental theorem of calculus]} \\
&= -\log S(y) && \text{[}S(0) = 1\text{, so } \log S(0) = 0\text{]} \\
\implies \quad S(y) &= e^{-\Theta(y)} && \text{[exponentiate]} \qquad (2) \\
f(y) &= \theta(y)\, e^{-\Theta(y)} && \text{[}f = \theta S \text{ by (1)]} \qquad (3)
\end{aligned}$$

- Θ: integrated hazard.
- θ, Θ, S, f: four descriptions of the same distribution.

1.3 Deeper look: survival = surviving many small coin flips

→ *Why now:* (2) came out of calculus. Here is what it *means*, and it shows the key fact that every hazard is about a selected group.

Split $[0, y]$ into n intervals of length $\Delta = y/n$, starting at $u_1, \ldots, u_n$:

$$\begin{aligned}
S(y) &= \prod_{k=1}^{n} \Pr\big(\text{survive interval } k \mid \text{survived to } u_k\big) && \text{[multiplication rule, repeated]} \\
&= \prod_{k=1}^{n} \big(1 - \theta(u_k)\Delta + o(\Delta)\big) && \text{[definition of the hazard]} \\
\log S(y) &= \sum_{k=1}^{n} \log\big(1 - \theta(u_k)\Delta + o(\Delta)\big) && \text{[log of a product]} \\
&= -\sum_{k=1}^{n} \theta(u_k)\Delta + o(1) && \text{[}\log(1-z) = -z + O(z^2)\text{]} \\
&\xrightarrow{n\to\infty} -\int_0^y \theta(u)\, du = -\Theta(y) && \text{[Riemann sum]}
\end{aligned}$$

- Being still in the state at y = winning a sequence of coin flips, each with exit probability $\theta(u_k)\Delta$.
- Each coin flip is made *only by those who won all earlier ones.*
  - So a hazard always describes a **selected group**: the survivors.
  - With heterogeneous people this selection changes who is in the group (Section 2).
- Given the hazard, *when* the exit happens is pure chance (the coin flips). Section 7 exploits this for treatment timing.

1.4 Example: constant hazard

→ *Why now:* the simplest benchmark; it defines "duration dependence", which Section 2 will show is hard to read from data.

$$\begin{aligned}
&\theta(y) = c \quad \Longrightarrow \quad \Theta(y) = cy, \qquad S(y) = e^{-cy} && \text{[integrate; then (2)]} \\
&\Pr(Y > y + u \mid Y > y) = \frac{S(y+u)}{S(y)} && \text{[conditional probability; } \{Y > y+u\} \subset \{Y > y\}\text{]} \\
&\qquad\qquad = e^{-cu} && \text{[plug in } S\text{]}
\end{aligned}$$

- Does not depend on y: **memoryless**.
- θ decreasing in y: **negative duration dependence**. Increasing: **positive**.

Takeaway → next: with one homogeneous population, the hazard is directly read from data via $\theta = f/S$. Next: what if people differ in ways we do not observe?

2 Unobserved heterogeneity and dynamic selection

→ *Why now:* workers differ in unobserved ways (e.g. motivation). The paper compares treated and not-yet-treated workers *among those still unemployed.* So we must know what heterogeneity does to the group that is still there.

Model

- Individuals differ by an unobserved type $V \geq 0$ (motivation, skills, ...), fixed during the spell.
- G: distribution of V in the population, $G(v) = \Pr(V \leq v)$.

  - Notation: $dG(v)$ = "probability that V is (around) v": $\Pr(V = v)$ if discrete, $g(v)\,dv$ if continuous.
  - So $\int h(v)\,dG(v) = \mathbb{E}[h(V)]$.
- Section 1 applies to *each type*: $\theta(y \mid V) = f(y \mid V)/S(y \mid V)$.
- **Assumption:** this conditional hazard has the form

$$\theta(y \mid V) = \lambda(y)\,V.$$

  - $\lambda(y) = \theta(y \mid V = 1) = f(y \mid V = 1)/S(y \mid V = 1)$: the **baseline**, the hazard of the $V = 1$ type.
  - Common time profile for all; V scales its level.
  - "$V = 1$ type" needs a normalization, since (λ, V) and $(c\lambda, V/c)$ give the same hazard.
- Integrated baseline $\Lambda(y) := \int_0^y \lambda$. So $\Theta(y \mid V) = \Lambda(y)V$ and, by (2), $\Pr(Y > y \mid V) = e^{-\Lambda(y)V}$.
- The **observed** hazard $\bar{\theta} = \bar{f}/\bar{S}$ (Section 1 applied to everyone pooled) is *not* assumed. It is derived below.

2.1 Observed survival and hazard

→ *Why now:* we only observe everyone pooled, not each type. Compute what the observed hazard looks like.

$$\begin{aligned}
\bar{S}(y) := \Pr(Y > y) &= \mathbb{E}\big[\Pr(Y > y \mid V)\big] && \text{[law of iterated expectations (LIE)]} \\
&= \mathbb{E}\big[e^{-\Lambda(y)V}\big] && [S(y \mid V) = e^{-\Lambda(y)V}] \\
\bar{\theta}(y) &= -\frac{\bar{S}'(y)}{\bar{S}(y)} && \text{[by (1) and } S' = -f] \\
&= \frac{\mathbb{E}\big[\lambda(y)Ve^{-\Lambda(y)V}\big]}{\mathbb{E}\big[e^{-\Lambda(y)V}\big]} && \text{[differentiate under } \mathbb{E};\ \Lambda' = \lambda] \\
&= \lambda(y)\,\frac{\mathbb{E}\big[Ve^{-\Lambda(y)V}\big]}{\mathbb{E}\big[e^{-\Lambda(y)V}\big]} && \text{[linearity: } \lambda(y) \text{ is not random]} \qquad (4)
\end{aligned}$$

- The ratio is an average of V, but with each type weighted by $e^{-\Lambda(y)V}$, its survival probability.
- Next: show this is the average type *among those still unemployed.*

2.2 Who survives: the type distribution among survivors

→ *Why now:* interpret the ratio in (4): it is about *who is left.*

Step 1: the distribution of types among survivors

$$\begin{aligned}
dG(v \mid Y > y) &= \frac{\Pr(Y > y \mid V = v)\,dG(v)}{\Pr(Y > y)} && \text{[Bayes' rule]} \\
&= \frac{S(y \mid v)\,dG(v)}{\bar{S}(y)} && \text{[definitions of } S(y \mid v),\ \bar{S}(y)] \\
&= \frac{e^{-\Lambda(y)v}\,dG(v)}{\mathbb{E}[e^{-\Lambda(y)V}]} && \text{[plug in } S(y \mid v) \text{ and } \bar{S}(y) \text{ from 2.1}] \qquad (5)
\end{aligned}$$

- $dG(v \mid Y > y)$: probability that a person *still unemployed at* y has type (around) v, i.e. the share of type v among survivors.

- Bayes flips the conditioning:
  - $S(y \mid v) = \Pr(\text{still there} \mid \text{type } v)$: given the type, the chance to survive;
  - $dG(v \mid Y > y) = \Pr(\text{type } v \mid \text{still there})$: given survival, the chance of the type.
- New share = old share $dG(v)$ × survival $S(y \mid v)$, renormalized.
- Bayes reading: prior $dG(v)$ (type mix at entry), likelihood $S(y \mid v)$ of the event "still unemployed", posterior $dG(v \mid Y > y)$. Staying unemployed is evidence of a low v.
- V is unobserved, so this cannot be counted in data; it follows from G and Λ once identified (Section 3).
- Example ($\lambda = 1$, types 0.3 and 1.5, half each, $y = 2$): $S(2 \mid v) \approx (0.55, 0.05)$, so the shares become $\approx (0.92, 0.08)$.

Step 2: the mean type among survivors

$$\begin{aligned}
\mathbb{E}[V \mid Y > y] &= \int v \, dG(v \mid Y > y) && \text{[conditional mean: average } v \text{ using the survivor distribution]} \\
&= \int v \, \frac{e^{-\Lambda(y)v} \, dG(v)}{\mathbb{E}[e^{-\Lambda(y)V}]} && \text{[plug in (5)]} \\
&= \frac{\int v \, e^{-\Lambda(y)v} \, dG(v)}{\mathbb{E}[e^{-\Lambda(y)V}]} && \text{[denominator does not depend on } v\text{: take it out]} \\
&= \frac{\mathbb{E}[V e^{-\Lambda(y)V}]}{\mathbb{E}[e^{-\Lambda(y)V}]} && [\textstyle\int h(v) \, dG(v) = \mathbb{E}[h(V)]] \qquad (6)
\end{aligned}$$

Step 3: back to the observed hazard

$$\bar{\theta}(y) = \lambda(y) \, \mathbb{E}[V \mid Y > y] \qquad \text{[(6) is the ratio in (4)]} \qquad (7)$$

- Example: $\mathbb{E}[V \mid Y > 2] \approx 0.92 \cdot 0.3 + 0.08 \cdot 1.5 \approx 0.40$, versus $\mathbb{E}[V] = 0.9$. So $\bar{\theta}(2) \approx 0.40$ vs. $\bar{\theta}(0) = 0.9$.
- The observed hazard averages over **survivors only**: $\lambda(y)\mathbb{E}[V \mid Y > y]$, not $\lambda(y)\mathbb{E}[V]$.

2.3 A name for the object that keeps appearing: the Laplace transform

→ *Why now:* every formula above contains $\mathbb{E}[e^{-aV}]$ for some a. Naming it and learning its properties is what makes all identification arguments below work.

Definition

$$\mathcal{L}(a) := \mathbb{E}\big[e^{-aV}\big] = \int_0^\infty e^{-av} \, dG(v), \qquad a \geq 0. \qquad (8)$$

- Reading in our model: $a = \Lambda(y)$ is the accumulated baseline hazard.
  - e^{-av}: probability that a type-v person is still unemployed.
  - $\mathcal{L}(a)$: share of everyone still unemployed. Indeed $\bar{S}(y) = \mathcal{L}(\Lambda(y))$.
- Same family as the moment generating function: $\mathcal{L}(a) = M_V(-a)$, $M_V(u) := \mathbb{E}[e^{uV}]$.
  - Always exists, since $V \geq 0$ gives $0 < e^{-aV} \leq 1$.

Properties

$$\begin{aligned}
\mathcal{L}(0) &= \mathbb{E}[1] = 1 && \text{[definition]}\\
\mathcal{L}'(a) &= -\mathbb{E}\big[Ve^{-aV}\big] < 0 && \text{[differentiate under } \mathbb{E}\text{; } < 0 \text{ if } \Pr(V>0)>0\text{]}\\
\mathcal{L}^{(k)}(a) &= (-1)^k\,\mathbb{E}\big[V^k e^{-aV}\big] && \text{[repeat } k \text{ times]}\\
-\mathcal{L}'(0) &= \mathbb{E}[V] && \text{[set } a=0\text{: derivatives at 0 give moments]}\\
-\frac{\mathcal{L}'(a)}{\mathcal{L}(a)} &= \frac{\mathbb{E}[Ve^{-aV}]}{\mathbb{E}[e^{-aV}]} = \mathbb{E}[V \mid \text{survived to } a] && \text{[by (6)]} \qquad (9)\\
\implies \quad \bar{\theta}(y) &= \lambda(y)\Big(-\frac{\mathcal{L}'}{\mathcal{L}}\Big)\big(\Lambda(y)\big) && \text{[by (7)]}
\end{aligned}$$

Two facts we state without proof:

- **Uniqueness:** $\mathcal{L}$ determines G. Two distributions with the same $\mathcal{L}$ are equal.
- **Analyticity:** for $a > 0$, $\mathcal{L}$ equals its Taylor series.
  - So knowing $\mathcal{L}$ on *any small interval* pins it down on all of $(0, \infty)$.
  - Like a polynomial: a short stretch of its graph fixes the whole graph.

Why this matters for identification

- Survival data are always $\mathcal{L}$ evaluated at some index (here $\Lambda(y)$).
- Strategy throughout: learn the index; learn $\mathcal{L}$ on an interval; then analyticity + uniqueness give G.

2.4 The mean type among survivors falls

→ *Why now:* show the *direction* of the composition effect in (7).

$$\begin{aligned}
\mathbb{E}[V \mid Y > y] &= -\frac{d}{da}\log\mathcal{L}(a)\Big|_{a=\Lambda(y)} && \text{[(9); chain rule: } (\log\mathcal{L})' = \mathcal{L}'/\mathcal{L}\text{]}\\
\frac{d^2}{da^2}\log\mathcal{L}(a) &= \frac{\mathcal{L}''}{\mathcal{L}} - \Big(\frac{\mathcal{L}'}{\mathcal{L}}\Big)^2 && \text{[quotient rule]}\\
&= \mathbb{E}[V^2 \mid \cdot] - \mathbb{E}[V \mid \cdot]^2 = \mathrm{Var}(V \mid \text{survived to } a) && \text{[properties above: survivor moments]}\\
\frac{d}{dy}\mathbb{E}[V \mid Y > y] &= -\lambda(y)\ \mathrm{Var}(V \mid Y > y)\ \le\ 0 && \text{[chain rule: } \Lambda' = \lambda\text{]}
\end{aligned}$$

- The observed hazard falls *even if no individual hazard falls.* This is **dynamic selection**.
- Hence two explanations for a falling exit rate:
  - true duration dependence: $\lambda \downarrow$ (e.g. skill loss: a *trap*);
  - composition: $\mathbb{E}[V \mid Y > y] \downarrow$ (sorting on permanent heterogeneity).
- People still unemployed at month 6 are a selected group. Every comparison made at month 6 inherits this.

2.5 Sorting and duration dependence cannot be separated

→ *Why now:* by (7) the observed hazard is $\lambda(y) \cdot \mathbb{E}[V \mid Y > y]$: duration dependence times sorting. Can data on Y split the two?

The identification question

- Data: the observed survival $\bar{S}(y)$ for all y (equivalently $\bar{\theta}(y) = -\frac{d}{dy}\log \bar{S}(y)$). That is all one duration gives us.
- Unknowns: the baseline Λ (duration dependence) and the type distribution G (sorting).
- The model links them: $\bar{S}(y) = \mathcal{L}_G\big(\Lambda(y)\big)$ (Section 2.3), where $\mathcal{L}_G$ is the Laplace transform of G.
- Not identified $\iff$ a *different* pair $(\tilde{\Lambda}, \tilde{G})$ gives the *same* $\bar{S}$.

Answer: no. For *any* $\tilde{G}$ there is a $\tilde{\Lambda}$ that reproduces $\bar{S}$ exactly (proof in Appendix A).

Why one can absorb the other

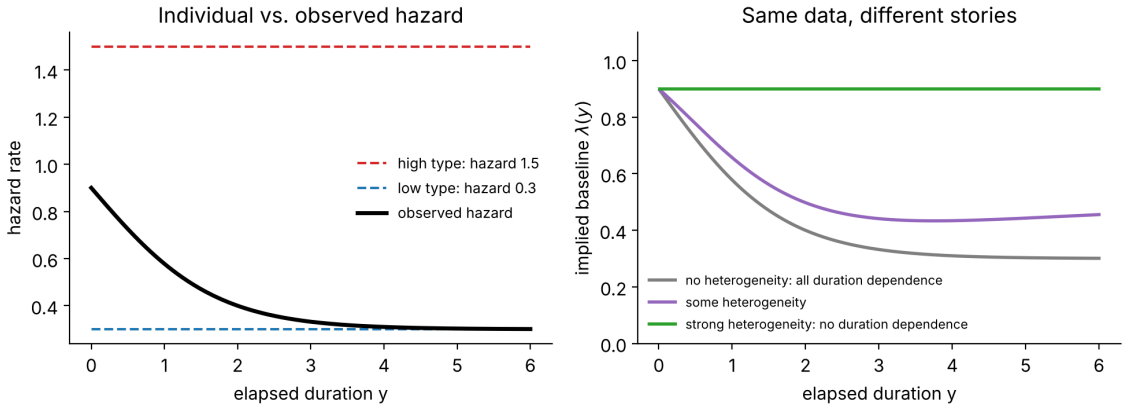
- Both act on the **same single dimension**, elapsed duration y:

$$\bar{\theta}(y) = \underbrace{\lambda(y)}_{\text{duration dependence}} \cdot \underbrace{\mathbb{E}[V \mid Y > y]}_{\text{sorting}}.$$

- Data give one number per y ($\bar{\theta}(y)$); the model has two unknown functions of y.
- λ is unrestricted, so whatever sorting path a G implies, λ can be bent to absorb the rest:
  - no heterogeneity: sorting is flat, so $\lambda = \bar{\theta}$ and the whole decline is duration dependence;
  - more heterogeneity: sorting falls faster (Section 2.4), so λ falls less.
- With one duration, **trap vs. sorting is a pure assumption**.

What would separate them: variation that moves *one* factor but not the other.

- **Groups at different stages of sorting at the same duration**, sharing the same λ.
  - Covariates that shift the hazard proportionally: the MPH model, Section 3.
- **The same person observed twice**, i.e. same V with the duration restarting.
  - Multiple spells, Section 9.
- **Restrictions on one factor**, e.g. a parametric λ (constant $\lambda \Rightarrow$ all decline is sorting) or a parametric G.
  - These identify by assumption, not by data.



Left: two types with constant hazards 0.3 and 1.5 (equal shares); the observed hazard falls from 0.9 to 0.3. Right: the same observed survival, fitted exactly by three type distributions (all with $\mathbb{E}V = 1$); each implies a different baseline $\tilde{\lambda}$.

Takeaway → next: the observed hazard mixes λ (true duration dependence) and G (who is left). Next: can data separate them? Only with structure, which is the MPH model.

3 The mixed proportional hazard (MPH) model

→ *Why now:* Section 2 showed that λ and G are confounded in the observed hazard. Here: what structure lets data separate them? This is the one-duration version of the argument the paper runs with two durations (Section 4).

Model

- $\theta(y \mid x, V) = \lambda(y)\,\varphi(x)\,V$, with $V \sim G$ and $V \perp X$.
- Three multiplicative pieces: duration (λ), observed (φ), unobserved (V).
- Normalize $\varphi(x^*) = 1$ and $\Lambda(y^*) = 1$ for some chosen x^*, y^*.
  - Reason: (λ, φ, V) and $(c\lambda, \varphi, V/c)$ give the same hazard.

3.1 Survival through the Laplace transform

→ *Why now:* write what the data reveal, $S(y \mid x)$ (equivalently its hazard $\theta(y \mid x)$), in terms of the unknowns (λ, φ, G).

$$\begin{aligned} S(y \mid x) &= \mathbb{E}\big[e^{-\varphi(x)\Lambda(y)V}\big] && \text{[LIE over } V\text{; (2); } V \perp X\text{]} \\ &= \mathcal{L}\big(\varphi(x)\Lambda(y)\big) && \text{[definition (8)]} \qquad (10) \\ \theta(y \mid x) &= \lambda(y)\varphi(x)\Big(-\frac{\mathcal{L}'}{\mathcal{L}}\Big)\big(\varphi(x)\Lambda(y)\big) && \text{[}-\tfrac{d}{dy}\log \text{ of (10); chain rule]} \qquad (11) \end{aligned}$$

- Data: $S(y \mid x)$ for all (y, x). The new information relative to Section 2.5 is the x dimension.
  - $\theta(y \mid x)$ is *not* extra information. Section 1: $S(y) = e^{-\Theta(y)}$ with $\Theta(y) = \int_0^y \theta$, and conversely $\theta = -\partial_y \log S$. Knowing the whole function $\theta \iff$ knowing the whole function S. We use whichever form makes an argument simpler.
- Data reveal one function $\mathcal{L}$ of a **single index** $\varphi(x)\Lambda(y)$.
- x and y both move the index: this is what the three steps below exploit.

3.2 Why covariates separate sorting from duration dependence

→ *Why now:* in Section 2.5 any G fit because Λ could be tailored to the one survival curve. With covariates, *one* Λ and *one* $\mathcal{L}$ must fit all groups at once: $S(y \mid x) = \mathcal{L}(\varphi(x)\Lambda(y))$ for every x. Show what this buys.

Let $m(a) := -\mathcal{L}'(a)/\mathcal{L}(a) = \mathbb{E}[V \mid \text{survived to } a]$. Compare groups x_1, x_2 with $\varphi_1 > \varphi_2$:

$$\begin{aligned} \frac{\theta(y \mid x_1)}{\theta(y \mid x_2)} &= \frac{\varphi_1}{\varphi_2} \cdot \frac{m\big(\varphi_1\Lambda(y)\big)}{m\big(\varphi_2\Lambda(y)\big)} && \text{[(11); } \lambda(y) \text{ cancels]} \qquad (12) \\ m'(a) &= -\operatorname{Var}(V \mid \text{survived to } a) && \text{[}m = -(\log\mathcal{L})'\text{; } (\log\mathcal{L})'' = \text{Var, Section 2.4]} \end{aligned}$$

- **No heterogeneity:** m is constant, so the ratio is φ_1/φ_2 at every y.
- **Heterogeneity** ($\operatorname{Var} V > 0$): m is strictly decreasing. For $y > 0$:

$$\begin{aligned} \varphi_1\Lambda(y) &> \varphi_2\Lambda(y) && \text{[}\varphi_1 > \varphi_2\text{, } \Lambda(y) > 0\text{]} \\ m\big(\varphi_1\Lambda(y)\big) &< m\big(\varphi_2\Lambda(y)\big) && \text{[}m \text{ strictly decreasing]} \\ \frac{\theta(y \mid x_1)}{\theta(y \mid x_2)} &= \frac{\varphi_1}{\varphi_2} \cdot \underbrace{\frac{m(\varphi_1\Lambda(y))}{m(\varphi_2\Lambda(y))}}_{<1} < \frac{\varphi_1}{\varphi_2} && \text{[(12)]} \end{aligned}$$

At $y = 0$: $\Lambda(0) = 0$, so the second factor is $m(0)/m(0) = 1$ and the ratio is φ_1/φ_2.

- The sign of $\theta(y \mid x_1)/\theta(y \mid x_2) - 1$ is *not* pinned down: with strong heterogeneity the ratio can fall below 1 (hazards cross). What we know, and what identification uses, is the drift *below* φ_1/φ_2.
- Why: at the same duration, the high-φ group is *further along the index*. Its high-V members have left faster, so its selection is more advanced.
- Duration dependence $\lambda(y)$ hits both groups equally and cancels in (12). Sorting does not cancel.
  - So the drift of the ratio away from φ_1/φ_2 measures sorting, separately from λ.
  - This is the separation that was impossible with one curve.
- The MPH assumption reads: *absent selection, covariate effects would be constant over duration.*
- Caveat: if true covariate effects vary with duration (non-proportional), drift is no longer evidence of sorting.

3.3 Identification (Elbers & Ridder 1982) in three steps

→ *Why now:* recover φ, then G, then λ from $S(y \mid x)$. Each step uses one property of $\mathcal{L}$ from Section 2.3. Section 4 repeats these steps in two dimensions.

Step 1: φ from short durations. Assume $\mathbb{E}V < \infty$.

$$\frac{\theta(y \mid x)}{\theta(y \mid x^*)} = \varphi(x)\,\frac{(-\mathcal{L}'/\mathcal{L})(\varphi(x)\Lambda(y))}{(-\mathcal{L}'/\mathcal{L})(\Lambda(y))} \qquad \text{[by (11); } \lambda(y) \text{ cancels; } \varphi(x^*) = 1\text{]}$$

$$\xrightarrow{y \downarrow 0} \varphi(x)\,\frac{-\mathcal{L}'(0)/\mathcal{L}(0)}{-\mathcal{L}'(0)/\mathcal{L}(0)} = \varphi(x) \qquad \text{[}\Lambda(y) \to 0\text{; continuity; } -\mathcal{L}'(0) = \mathbb{E}V < \infty\text{]}$$

- At $y \approx 0$ nobody has been selected out yet, so the ratio is free of selection.
- Uses short durations, where data are thin.

Step 2: G at the normalization date.

$$S(y^* \mid x) = \mathcal{L}\big(\varphi(x)\big) \qquad \text{[(10) with } \Lambda(y^*) = 1\text{]}$$

- $\varphi(x)$ known from step 1. Varying x traces $\mathcal{L}$ on an interval.
- Analytic $\Rightarrow$ $\mathcal{L}$ known on $(0, \infty)$ $\Rightarrow$ G known.

Step 3: baseline.

$$\Lambda(y) = \frac{\mathcal{L}^{-1}\big(S(y \mid x)\big)}{\varphi(x)} \qquad \text{[invert (10); } \mathcal{L} \text{ strictly decreasing]}$$

What was assumed: proportionality, $V \perp X$, $\mathbb{E}V < \infty$, variation in $\varphi(x)$.

Takeaway → next: with MPH structure and covariates, one duration reveals λ, φ and the full distribution of unobserved types. In the paper, *before* treatment, two durations compete: job vs. benefits reduction. Their link through unobserved types is exactly the selection problem. Next: extend the argument to two durations.

4 Two durations: competing risks

→ *Why now:* Section 3 learned the heterogeneity of *one* duration. The treatment brings a second one.

Where the second duration comes from

- The treatment is itself a duration: the benefits reduction is not assigned at a fixed date, it *arrives* at a random time S with its own hazard.
- So two durations run at once: Y (time to a job) and S (time to the reduction).
- Two phases:
  - **before S: competing risks**; whichever comes first ends it (this section);
  - **after S:** Y continues under treatment; the effect δ shows up here (Sections 7–8).

Why it is a hurdle

- **Censoring:** if the job comes first, we never see when the reduction would have come.
- **Endogeneity:** the same unobserved traits may drive both (low motivation $\Rightarrow$ slow job finding *and* early reduction), so V_Y and V_S are correlated. This *is* the selection problem.
- **Negative result** (Tsiatis, Section 4.5): competing-risks data alone cannot reveal that correlation.

Why modelling it is also the solution

- If S were exogenous ($V_S \perp V_Y$), "reduced yet?" would be an ordinary time-varying covariate in Y's hazard, and Section 3 would suffice.
- It is not, so we model S's own process: the paper's dynamic **selection equation** (as in Heckman's selection model). This buys:
  - **A clean window:** before S there is no causal effect (no anticipation, Section 5), so the competing-risks data reflect only the joint heterogeneity of (V_Y, V_S).
  - **The correlation becomes identifiable:** MPH plus covariates beat Tsiatis (Section 4.6), the two-duration version of Section 3.
  - **Instrument-like variation:** given (x, V), the timing of S still contains pure luck (Section 7).

This section studies the competing-risks part alone. No treatment effect is involved yet.

Set-up

- Two latent durations: Y (exit to work) and S (benefits reduction).
  - Here Y is the time to a job *if never treated* ($= Y^*(\infty)$ in Section 5).
- Only the **first** event is seen, plus which one it was.
  - If Y comes first: S is censored.
  - If S comes first: we do see a job time later, but it is a *treated* outcome, not this Y. It is used in Sections 6–8, not here.
  - So the pair of latent durations (Y, S) is never seen in full.
- MPH for each, with **correlated** unobservables $V = (V_Y, V_S) \sim G$, $V \perp X$:

$$\theta_Y(t \mid x, V) = \lambda_Y(t)\varphi_Y(x)V_Y, \qquad \theta_S(t \mid x, V) = \lambda_S(t)\varphi_S(x)V_S.$$

- **Assumption** (competing-risks version of the paper's Assumption 2, Section 7): $Y \perp S \mid (x, V)$.
  - Given observed x and unobserved type V, the luck in job timing is unrelated to the luck in reduction timing.
  - All dependence between Y and S runs through (x, V): selection on observables and unobservables.
  - Rules out shocks outside V that hit both (e.g. a bad week that lowers job finding *and* triggers a reduction).
- Bivariate Laplace transform: $\mathcal{L}(a, b) := \mathbb{E}[e^{-aV_Y - bV_S}]$; partials $\mathcal{L}_a, \mathcal{L}_b, \mathcal{L}_{ab}$.
  - Same properties as Section 2.3: derivatives give moments among survivors; uniqueness; analyticity.
- Integrated baselines with covariates (no V): $\Lambda_Y(t \mid x) := \varphi_Y(x)\Lambda_Y(t)$, $\Lambda_S(t \mid x) := \varphi_S(x)\Lambda_S(t)$.

4.1 Joint survival

→ *Why now:* the model's implication for the two latent durations together.

$$\begin{aligned}
\Pr(Y > y, S > s \mid x) &= \mathbb{E}\big[\Pr(Y > y, S > s \mid x, V)\big] && \text{[LIE over } V\text{; } V \perp X\text{]} \\
&= \mathbb{E}\big[\Pr(Y > y \mid x, V)\Pr(S > s \mid x, V)\big] && \text{[} Y \perp S \mid (x, V)\text{]} \\
&= \mathbb{E}\big[e^{-\varphi_Y(x)\Lambda_Y(y)V_Y}\, e^{-\varphi_S(x)\Lambda_S(s)V_S}\big] && \text{[(2) for each]} \\
&= \mathcal{L}\big(\varphi_Y(x)\Lambda_Y(y),\, \varphi_S(x)\Lambda_S(s)\big) && \text{[definition of } \mathcal{L}\text{]} \qquad (13)
\end{aligned}$$

4.2 What the data reveal

→ *Why now:* we never see both durations, only the first event. Derive which pieces of (13) are observable.

Probability that nothing has happened by t (the diagonal $y = s = t$):

$$\bar{S}_0(t \mid x) := \Pr(Y > t, S > t \mid x) = \mathcal{L}\big(\Lambda_Y(t \mid x), \Lambda_S(t \mid x)\big) \qquad \text{[(13) at } y = s = t\text{]} \qquad (14)$$

Cause-specific hazard of exit (rate of Y happening first, among those with nothing yet):

$$\begin{aligned}
\bar{\theta}_0^Y(t \mid x) &:= \lim_{dt \downarrow 0} \frac{\Pr(t \le Y < t + dt,\ S > t \mid x)}{dt\ \bar{S}_0(t \mid x)} && \text{[definition]} \\
&= \frac{-\partial_y \Pr(Y > y, S > t \mid x)\big|_{y=t}}{\bar{S}_0(t \mid x)} && \text{[derivative of a survival in } y\text{]} \\
&= \frac{-\mathcal{L}_a(\Lambda_Y(t \mid x), \Lambda_S(t \mid x))\, \varphi_Y(x)\lambda_Y(t)}{\mathcal{L}(\Lambda_Y(t \mid x), \Lambda_S(t \mid x))} && \text{[chain rule on (13); } \Lambda_Y' = \lambda_Y\text{]} \\
&= \lambda_Y(t)\varphi_Y(x)\,\mathbb{E}[V_Y \mid R_t, x] && \text{[}-\mathcal{L}_a/\mathcal{L} = \text{survivor mean, as in (6)]} \qquad (15)
\end{aligned}$$

where $R_t := \{Y > t, S > t\}$ is the **risk set**. Symmetrically, $\bar{\theta}_0^S(t \mid x) = \lambda_S(t)\varphi_S(x)\mathbb{E}[V_S \mid R_t, x]$.

- Notation: $\bar{S}_0, \bar{\theta}_0$ = observed survival and hazards *before treatment* (later $\bar{S}_1$: after). Superscript: which exit.
- Same structure as Section 1, split by cause: $\bar{\theta}_0^Y + \bar{\theta}_0^S = -\frac{d}{dt}\log \bar{S}_0$.
- $\bar{\theta}_0^Y$ is *cause-specific*: it conditions on $S > t$ too, so it is not the hazard of Y (unless $V_Y \perp V_S$).

4.3 Where the dependence of V_Y, V_S lives

→ *Why now:* the selection is the dependence of V_Y and V_S. Locate it inside the observable $\bar{S}_0$.

Write $w(v) \propto e^{-av_Y - bv_S}$ (survivor weights) and $\mathbb{E}_w$ for expectations under them.

$$\begin{aligned}
\partial_a \log \mathcal{L} &= \frac{\mathcal{L}_a}{\mathcal{L}} = -\mathbb{E}_w[V_Y] && \text{[chain rule; differentiate under } \mathbb{E}\text{]} \\
\partial_b\partial_a \log \mathcal{L} &= \frac{\mathcal{L}_{ab}}{\mathcal{L}} - \frac{\mathcal{L}_a\mathcal{L}_b}{\mathcal{L}^2} && \text{[quotient rule]} \\
&= \mathbb{E}_w[V_Y V_S] - \mathbb{E}_w[V_Y]\mathbb{E}_w[V_S] && \text{[}\mathcal{L}_{ab} = \mathbb{E}[V_Y V_S e^{\cdots}],\ \mathcal{L}_b/\mathcal{L} = -\mathbb{E}_w[V_S]\text{]} \\
&= \mathrm{Cov}(V_Y, V_S \mid R_t, x) && \text{[definition of covariance; at } (a,b) = (\Lambda_Y(t \mid x), \Lambda_S(t \mid x))\text{]} \qquad (16)
\end{aligned}$$

- $V_Y \perp V_S \iff \mathcal{L}(a,b) = \mathcal{L}_Y(a)\mathcal{L}_S(b) \iff \log \bar{S}_0$ additively separable in the two indices.
- Dependence (selection) = an **interaction** of the two indices in $\log \bar{S}_0$.

4.4 Deeper look: how selection shows up in competing risks

→ *Why now:* (16) is abstract. Translate it into something observable: how a shift in treatment intensity changes the exit rate of those *not* treated.

Hold the exit index a fixed and raise the treatment index b (a group where reductions happen more often):

$$\begin{aligned}
\bar{\theta}_0^Y &= \lambda_Y \varphi_Y \, \mathbb{E}_w[V_Y], \qquad w \propto e^{-aV_Y - bV_S} && \text{[(15)]} \\
\frac{\partial}{\partial b}\mathbb{E}_w[V_Y] &= -\mathbb{E}_w[V_Y V_S] + \mathbb{E}_w[V_Y]\mathbb{E}_w[V_S] && \text{[quotient rule, as in Section 2.4]} \\
&= -\operatorname{Cov}(V_Y, V_S \mid R_t, x) && \text{[definition of covariance]} \\
\implies \quad \frac{\partial \log \bar{\theta}_0^Y}{\partial b} &= -\frac{\operatorname{Cov}(V_Y, V_S \mid R_t, x)}{\mathbb{E}[V_Y \mid R_t, x]} && \text{[log derivative; } \lambda_Y\varphi_Y \text{ do not depend on } b\text{]} \quad (17)
\end{aligned}$$

- Those in the risk set are *untreated*. By no anticipation, treatment intensity has no causal effect on their exit rate.
- So any response of their exit rate to b is **pure selection**:
  - more treatment removes high-V_S people from the untreated pool;
  - if high V_S goes with high V_Y, the untreated survivors are left with lower V_Y, so their exit rate falls.
- No selection ($\operatorname{Cov} = 0$): shifting treatment intensity leaves the untreated exit rate unchanged.
- This is the IV-like role of covariates:
  - no variable needs to be excluded from the exit hazard;
  - but we must move b *holding a fixed*, i.e. x must shift the two indices in different directions (A3).
- Treatment acts like a competing exit: it removes people from the untreated group, just as job-finding removes people in Section 2.

Assumption 3 (paper's): non-proportional covariate variation.

$$\big\{\big(\varphi_Y(x), \varphi_S(x)\big) : x \in \mathcal{X}\big\} \text{ contains a non-empty open set of } \mathbb{R}^2.$$

- Meaning: varying x moves the two indices in genuinely different directions, so b can move with a fixed.
- Example: $\varphi_j(x) = e^{x'\beta_j}$, $x = (x_1, x_2)$ continuous. Move x in direction $d = (\beta_{Y2}, -\beta_{Y1})$:

$$\begin{aligned}
d'\beta_Y &= \beta_{Y2}\beta_{Y1} - \beta_{Y1}\beta_{Y2} = 0 && [a \text{ fixed}] \\
d'\beta_S &= \beta_{Y2}\beta_{S1} - \beta_{Y1}\beta_{S2} \neq 0 && [\text{iff } (\beta_Y \ \beta_S) \text{ full rank: } b \text{ moves}]
\end{aligned}$$

vs. an exclusion restriction

- Exclusion: a *named* variable has zero effect on Y's hazard. Assumed, and hard to defend: agents react to what changes their treatment risk.
- Assumption 3: no variable is excluded. Some *direction* d is excluded from Y's index, and it is **learned** from data (φ_Y, φ_S are identified first, Section 4.6, step 1).
- The price moves elsewhere: x must affect each hazard only through one multiplicative index, constant over duration (MPH).

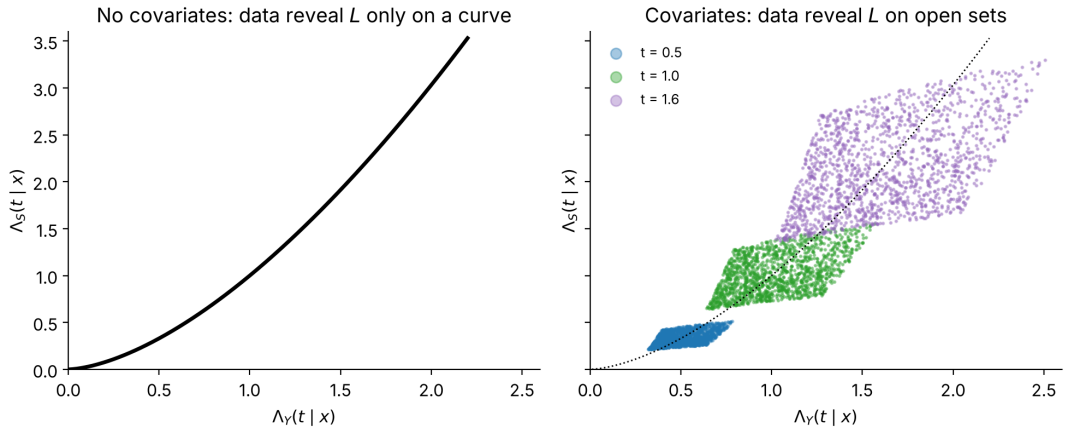
4.5 The problem: Tsiatis (1975)

→ *Why now:* the dependence sits in $\mathcal{L}$, but do the data reveal enough of $\mathcal{L}$? Without structure, no.

- Without structure, any competing-risks data can be generated by *independent* Y, S.
- Geometry, without covariates:
  - the data reveal $\mathcal{L}$ only at points $(\Lambda_Y(t), \Lambda_S(t))$, i.e. along a **curve**;
  - a curve cannot pin down a function of two variables, so the dependence is not identified.
- With covariates moving φ_Y and φ_S separately: points $(\Lambda_Y(t \mid x), \Lambda_S(t \mid x))$ fill **open sets**.

How to read the figure

- A point is a *location* (a, b): exit index a and treatment index b.
- Over each location sits the number $\mathcal{L}(a, b)$: the probability that *nothing has happened yet* (no job, no reduction) at those index values. Think of it as a height, a third dimension not drawn.
- The shape of this surface is the unknown; the dependence of V_Y, V_S (selection) lives in it.
- Each plotted point is one group x at one duration t:
  - location: $\big(\varphi_Y(x)\Lambda_Y(t),\ \varphi_S(x)\Lambda_S(t)\big)$, from model pieces identified in Section 4.6 (steps 1 and 3);
  - height: $\bar{S}_0(t \mid x)$, the share of group x with neither event by t. This is **estimated from data**.
- So at plotted locations the height is known; elsewhere it is not. Laplace properties (analyticity, uniqueness) are what extend the known heights to the whole surface and then to G.
- Left (one group): as t grows, the location moves along one path $t \mapsto (\Lambda_Y(t), \Lambda_S(t))$.
  - A straight line if $\Lambda_S \propto \Lambda_Y$; a curve otherwise. Here $\Lambda_Y(t) = t$, $\Lambda_S(t) = t^{1.6}$, for illustration.
  - Either way it is **one-dimensional**: heights along one path cannot pin down a surface.
- Right (many groups): at a fixed t (one colour), groups spread over a **two-dimensional patch** because x moves the two indices differently (Assumption 3).
  - Heights on a patch pin down the whole surface (analyticity).
  - Dotted line: the left panel's path.



4.6 Identification (Prop. 2) in three steps

→ *Why now:* show that MPH + covariates fix the Tsiatis problem: the steps of Section 3.3, in two dimensions.

Assumptions. (A3) Assumption 3, Section 4.4. (A4) $\mathbb{E}V_Y, \mathbb{E}V_S < \infty$. Normalize $\varphi_Y(x^*) = \varphi_S(x^*) = 1$, $\Lambda_Y(t^*) = \Lambda_S(t^*) = 1$.

Step 1: φ_Y, φ_S**.**

$$\frac{\bar\theta_0^Y(t \mid x)}{\bar\theta_0^Y(t \mid x^*)} = \varphi_Y(x)\,\frac{\mathbb{E}[V_Y \mid R_t, x]}{\mathbb{E}[V_Y \mid R_t, x^*]} \qquad \text{[by (15); } \lambda_Y(t) \text{ cancels]}$$

$$\xrightarrow{t\downarrow 0} \varphi_Y(x)\,\frac{\mathbb{E}V_Y}{\mathbb{E}V_Y} = \varphi_Y(x) \qquad \text{[risk set} \to \text{everyone; } V \perp X\text{; (A4)]}$$

Same for φ_S with $\bar\theta_0^S$. Uses *which* event happened.

Step 2: G**, including the dependence.**

$$\bar S_0(t^* \mid x) = \mathcal{L}\big(\varphi_Y(x), \varphi_S(x)\big) \qquad \text{[(14) with } \Lambda_Y(t^*) = \Lambda_S(t^*) = 1]$$

- Known points fill an open set (A3) $\Rightarrow$ $\mathcal{L}$ known there.
- Analytic $\Rightarrow$ known on $(0,\infty)^2$ $\Rightarrow$ joint G known.

Step 3: baselines. For each t, $(\Lambda_Y(t), \Lambda_S(t))$ is the unique (a, b) with

$$\bar S_0(t \mid x) = \mathcal{L}\big(\varphi_Y(x)a,\ \varphi_S(x)b\big) \quad \text{for all } x. \qquad \text{[(14); } \mathcal{L}, \varphi\text{'s known]}$$

- Uniqueness: two solutions would give $\mathcal{L}(u,v) = \mathcal{L}(ku, mv)$ on $(0,\infty)^2$, $(k,m) \neq (1,1)$.
- Iterating, $\mathcal{L}$ would be constant, contradicting strict monotonicity.
- These three steps are our reconstruction; the paper cites Abbring & Van den Berg (2003, JRSSB).

Takeaway $\to$ **next:** the competing-risks data alone reveal $\lambda_Y, \lambda_S, \varphi_Y, \varphi_S$ and the joint G, i.e. how selected those still at risk are. Next: introduce the treatment and define its causal effect.

5 Adding a treatment

$\to$ *Why now:* so far no causal effect. We need to define "the effect of a benefits reduction at s", and say when the pre-treatment competing-risks part of Section 4 is untouched by the treatment.

5.1 Potential outcomes indexed by treatment time

$\to$ *Why now:* a dynamic treatment is a choice of *date*, so the causal objects are one outcome per date.

- Treatment = benefits reduction at date $s \in [0,\infty]$; $s = \infty$: never.
- One potential outcome per date: $Y^*(s)$, with hazard $\theta_{Y^*(s)}$, integrated hazard $\Theta_{Y^*(s)}$.
- Each $Y^*(s)$ is fully described by its hazard, so "the effect of s" = how the hazard changes with s.
- Actual treatment date S; observed outcome $Y = Y^*(S)$.

5.2 No anticipation (Assumption 1)

$\to$ *Why now:* for Section 4 to apply before treatment, the future treatment date must not affect the hazard before it.

$$\Theta_{Y^*(s)}(y) = \Theta_{Y^*(s')}(y) \qquad \text{for all } y \le \min\{s, s'\}.$$

Take $s' = \infty$ and differentiate in y:

$$\theta_{Y^*(s)}(y) = \theta_{Y^*(\infty)}(y) \qquad \text{for } y < s. \qquad \text{[derivative of both sides]}$$

- Before treatment, everyone has the **untreated** hazard, whatever their future date.
- Economic content: people do not act on knowledge of their own future reduction date.
- Violated if, e.g., the date is announced in advance and people search harder before it.
- Not ruled out: reacting to the *policy* (e.g. the rate at which reductions happen).
  - This reaction is in every $Y^*(s)$.
  - So the effect is "treated at s vs. later, under this policy", not "policy vs. no policy".
- Prop. 1: without more structure, data cannot distinguish (i) causal effects from selection, nor (ii) anticipation from no anticipation. Hence Assumption 1 is maintained, and selection must be modelled.

Takeaway → next: under no anticipation, the period before treatment is a pure competing-risks problem, as in Section 4. Next: what exactly do the data contain?

6 What we actually observe

→ *Why now:* identification means recovering the model from the data distribution. We need that distribution precisely, including what is censored.

The censoring

- S is seen only if $S < Y$: if the worker finds a job first, the reduction date is never seen.
- Paper's data (per x): $Q_S(y,s) := \Pr(Y > y, S > s, Y > S)$ and $Q_Y(y) := \Pr(Y > y, Y < S)$. Assume $\Pr(Y = S) = 0$.

Pre-treatment objects

$$\begin{aligned}
\bar{S}_0(t) &= \Pr(Y > t, S > t, Y > S) + \Pr(Y > t, S > t, Y < S) && \text{[partition by } Y \gtrless S;\ \Pr(Y=S)=0\text{]} \\
&= Q_S(t,t) + Q_Y(t) && [Y > t, Y < S \text{ implies } S > t] \\
\bar{\theta}_0^Y(t) &= \frac{-Q_Y'(t)}{\bar{S}_0(t)} && [-Q_Y' \text{: density of untreated exits}] \\
\bar{\theta}_0^S(t) &= \frac{-\partial_s Q_S(y,s)\big|_{y\downarrow t,\, s=t}}{\bar{S}_0(t)} && \text{[density of } S \text{ at } t\text{, still unemployed]}
\end{aligned}$$

Post-treatment object: survival of the cohort treated at s. For $y > s$:

$$\begin{aligned}
-\partial_s Q_S(y,s) &= \lim_{ds\downarrow 0} \frac{\Pr(Y > y,\ s \le S < s + ds,\ Y > S)}{ds} && \text{[derivative of } Q_S \text{ in } s\text{]} \\
&= \lim_{ds\downarrow 0} \frac{\Pr(Y > y,\ s \le S < s + ds)}{ds} && [Y > y > S \text{ implies } Y > S] \\
\bar{S}_1(y \mid s) &:= \Pr(Y > y \mid S = s,\, Y > s) && \text{[definition]} \\
&= \frac{\partial_s Q_S(y,s)}{\lim_{y'\downarrow s} \partial_s Q_S(y',s)} && \text{[cond. probability; } ds \text{ cancels]} \qquad (18)
\end{aligned}$$

- $(\bar{S}_0, \bar{\theta}_0^Y, \bar{\theta}_0^S)$ = competing-risks data, as in Section 4.
- $\bar{S}_1$ = the only place where the treatment effect can show up.

Takeaway → next: pre-treatment data identify the competing-risks model (Section 4). Post-treatment data $\bar{S}_1$ must deliver δ. Next: the model linking both.

7 The model of the paper (Model 1a)

→ *Why now:* combine the pieces: the competing-risks MPH model of Section 4 before treatment, and a causal effect δ after it.

Assumption 2 (selection on observables and unobservables): $Y^*(\cdot) \perp S \mid (X, V)$, with V time-invariant.

Model 1a

$$\theta_Y(t \mid s, x, V) = \begin{cases} \lambda_Y(t)\varphi_Y(x)V_Y & t \le s \\ \lambda_Y(t)\varphi_Y(x)V_Y\,\delta(t \mid s, x) & t > s \end{cases}$$
$$\theta_S(t \mid x, V) = \lambda_S(t)\varphi_S(x)V_S, \qquad (V_Y, V_S) \sim G, \quad V \perp X.$$

- δ: **the treatment effect**, a multiplier on the exit rate after treatment.
  - May depend on t, s (hence on time since treatment $t - s$), and x; not on V.
- No anticipation is built in: δ only acts for $t > s$.
- Selection on unobservables: V_Y and V_S may be dependent in any way.
- The same x in both hazards: **no exclusion restriction**.
- $\lambda_Y, \lambda_S, \varphi_Y, \varphi_S, G, \delta$ all nonparametric.
- Plus (A3), (A4) from Section 4.

Integrated outcome hazard, with $D(y \mid s, x) := \int_s^y \lambda_Y(\tau)\delta(\tau \mid s, x)\,d\tau$:

$$\begin{aligned} \Theta_Y(y \mid s, x, V) &= \int_0^s \lambda_Y\varphi_Y V_Y\,d\tau + \int_s^y \lambda_Y\varphi_Y V_Y\delta\,d\tau && \text{[split the integral at } s\text{]} \\ &= \varphi_Y(x)V_Y\big[\Lambda_Y(s) + D(y \mid s, x)\big] \quad (y > s) && \text{[linearity of the integral]} \end{aligned} \tag{19}$$

Which part of the treatment time is pure chance?

→ *Why now:* for causal inference we need variation in treatment that is unrelated to the outcome. In IV it comes from an excluded variable; here there is none. So: given (x, V), is anything left in S that is pure luck?

Trick: transform S by its own integrated hazard, $E_S := \Theta_S(S \mid x, V) = \varphi_S(x)V_S\Lambda_S(S)$. Fix (x, V); for $e \ge 0$:

$$\begin{aligned} \Pr(E_S > e \mid x, V) &= \Pr\big(\Theta_S(S \mid x, V) > e \mid x, V\big) && \text{[definition of } E_S\text{]} \\ &= \Pr\big(S > \Theta_S^{-1}(e \mid x, V) \mid x, V\big) && [\Theta_S \text{ strictly increasing in } S] \\ &= \exp\big(-\Theta_S(\Theta_S^{-1}(e \mid x, V) \mid x, V)\big) && [(2), \text{ conditional on } (x, V)] \\ &= e^{-e} && [\Theta_S(\Theta_S^{-1}(e)) = e] \end{aligned}$$

- The right side does not depend on (x, V). So $E_S \sim \text{Exp}(1)$ and $E_S \perp (x, V)$.
- (Same as $F(S) \sim U(0, 1)$ for any continuous S: indeed $E_S = -\log \Pr(S > s \mid x, V)\big|_{s=S}$.)

Invert the definition of E_S:

$$S = \Lambda_S^{-1}\Big(\frac{E_S}{\varphi_S(x)\,V_S}\Big) \qquad \text{[solve } \varphi_S V_S \Lambda_S(S) = E_S \text{ for } S\text{]}$$

- The treatment time splits into a **systematic part** (x, V_S) and **pure luck** E_S.
- V_S is the source of selection (it may be correlated with V_Y).

- E_S moves S around, and by Assumption 2 it does not enter Y's hazard.
  - It plays the role of an instrument. But what is excluded is an unobserved **shock**, not an observed variable.
  - This is the paper's footnote 15: in each instant, treatment is a coin flip with probability $\theta_S\,dt$ (Section 1.3).
- Contrast with a static threshold model $D = \mathbf{1}\{\mu(x) > U\}$: given (x, U), D is deterministic, so no luck is left.

Takeaway → next: the model has a random component in timing (like an instrument) and selection through V (to be learned from the competing-risks part). Next: show δ is identified.

8 Identification, step by step

→ *Why now:* everything is in place. Plan:

- A: learn the selection (G) from pre-treatment data;
- B: describe who is treated at s;
- C: the simplest case, the moment of treatment;
- D: the general case.

8.1 Step A: the pre-treatment data are the competing-risks model of Section 4

→ *Why now:* check that the pre-treatment data are exactly the competing-risks model of Section 4, so its results apply.

$$\begin{aligned}
\bar{S}_0(t \mid x) &= \mathbb{E}\big[\Pr(Y > t, S > t \mid x, V)\big] && \text{[LIE; } V \perp X\text{]} \\
&= \mathbb{E}\big[e^{-\varphi_Y(x)\Lambda_Y(t)V_Y}\, e^{-\varphi_S(x)\Lambda_S(t)V_S}\big] && \text{[Assumption 2; before } S\text{, } \delta \text{ absent (no anticipation)]} \\
&= \mathcal{L}\big(\Lambda_Y(t \mid x), \Lambda_S(t \mid x)\big) && \text{[as in (14)]}
\end{aligned}$$

- Same for $\bar{\theta}_0^Y, \bar{\theta}_0^S$. So Prop. 2 (Section 4.6) applies: $\varphi_Y, \varphi_S, \Lambda_Y, \Lambda_S, G$ are identified.
- This uses no post-treatment data.

8.2 Step B: who gets treated at s

→ *Why now:* the effect is estimated by following the cohort treated at s. We need its type distribution.

Risk set at s:

$$\begin{aligned}
dG_{R_s}(v) &= \frac{\Pr(R_s \mid x, v)\, dG(v)}{\Pr(R_s \mid x)} && \text{[Bayes' rule]} \\
&\propto e^{-a_s v_Y - b_s v_S}\, dG(v) && \text{[as in Step A]}
\end{aligned}$$

Cohort treated at s (in the risk set and treated in $[s, s + ds)$):

$$\begin{aligned}
dG_s(v) &\propto \Pr(s \le S < s + ds \mid R_s, x, v)\, dG_{R_s}(v) && \text{[Bayes' rule]} \\
&\propto \lambda_S(s)\varphi_S(x)v_S\, ds\, dG_{R_s}(v) && \text{[definition of the hazard } \theta_S\text{]} \\
&\propto v_S\, e^{-a_s v_Y - b_s v_S}\, dG(v) && \text{[drop factors not depending on } v\text{]} \qquad (20)
\end{aligned}$$

- Treated cohort = risk set **reweighted by the treatment intensity** v_S.
- Fully known after Step A.

8.3 Step C: the selection problem at the moment of treatment

→ *Why now:* the cleanest comparison: right at s, compare the just-treated with those not yet treated. This isolates the selection problem in one factor.

Exit hazard of the just-treated vs. the not-yet-treated, at s:

$$\begin{aligned} \text{just treated:} \quad & \lambda_Y(s)\varphi_Y(x)\delta(s \mid s,x)\,\mathbb{E}_{G_s}[V_Y] && \text{[Model 1a, averaged over (20)]} \\ \text{not yet treated:} \quad & \lambda_Y(s)\varphi_Y(x)\,\mathbb{E}_{R_s}[V_Y] && \text{[(15)]} \end{aligned}$$

$$\begin{aligned} \mathbb{E}_{G_s}[V_Y] &= \frac{\mathbb{E}_{R_s}[V_Y V_S]}{\mathbb{E}_{R_s}[V_S]} && \text{[(20): } G_s = G_{R_s} \text{ reweighted by } v_S\text{]} \\ &= \mathbb{E}_{R_s}[V_Y] + \frac{\mathrm{Cov}_{R_s}(V_Y, V_S)}{\mathbb{E}_{R_s}[V_S]} && \text{[}\mathbb{E}[XY] = \mathbb{E}X\mathbb{E}Y + \mathrm{Cov}(X,Y)\text{]} \\ \frac{\text{just treated}}{\text{not yet treated}} &= \delta(s \mid s,x)\Big(1 + \frac{\mathrm{Cov}_{R_s}(V_Y, V_S)}{\mathbb{E}_{R_s}[V_Y]\mathbb{E}_{R_s}[V_S]}\Big) && \text{[divide; } \lambda_Y\varphi_Y \text{ cancel]} \\ &= \delta(s \mid s,x) \cdot \frac{\mathcal{L}\,\mathcal{L}_{ab}}{\mathcal{L}_a\,\mathcal{L}_b}\Big|_{(\Lambda_Y(s|x),\Lambda_S(s|x))} && \text{[survivor moments from Section 4.3]} \end{aligned}$$

- Observed ratio = causal effect × **selection factor**.
- Selection factor = 1 iff V_Y, V_S uncorrelated in the risk set: then timing is as good as random at s.
- The selection factor is known from Step A, so $\delta(s \mid s,x)$ is identified.

8.4 Step D: beyond the moment of treatment (Prop. 3)

→ *Why now:* Step C only gives δ at $t = s$. After s the treated cohort keeps being reshaped by exits, so we need its full survival curve.

For $y > s$, let $g(y, s \mid x)$ be the density of "treated at s and still unemployed at y", and $D := D(y \mid s,x)$:

$$\begin{aligned} g(y,s \mid x) &= \mathbb{E}\big[f_S(s \mid x,V)\ \Pr(Y > y \mid S = s,x,V)\big] && \text{[LIE; Assumption 2]} \\ &= \mathbb{E}\big[\lambda_S(s)\varphi_S V_S e^{-b_s V_S} \cdot e^{-\varphi_Y V_Y[\Lambda_Y(s)+D]}\big] && \text{[(3); (2) with (19)]} \\ &= \lambda_S(s)\varphi_S\,\mathbb{E}\big[V_S\, e^{-(\Lambda_Y(s|x)+\varphi_Y D)V_Y - b_s V_S}\big] && \text{[linearity; } \Lambda_Y(s \mid x) = \varphi_Y\Lambda_Y(s)\text{]} \\ &= -\lambda_S(s)\varphi_S(x)\,\mathcal{L}_b\big(\Lambda_Y(s \mid x) + \varphi_Y(x)D,\ \Lambda_S(s \mid x)\big) && \text{[}\mathcal{L}_b = -\mathbb{E}[V_S e^{-aV_Y - bV_S}]\text{]} \end{aligned}$$

Divide by the same expression at $y \downarrow s$ (where $D = 0$), as in (18):

$$\begin{aligned} \bar{S}_1(y \mid s,x) &= \frac{\mathcal{L}_b\big(\Lambda_Y(s \mid x) + \varphi_Y(x)D(y \mid s,x),\ \Lambda_S(s \mid x)\big)}{\mathcal{L}_b(\Lambda_Y(s \mid x), \Lambda_S(s \mid x))} && \text{[the } \lambda_S\varphi_S \text{ cancel]} \\ &= \mathbb{E}_{G_s}\big[e^{-\varphi_Y(x)V_Y D(y|s,x)}\big] && \text{[rewrite with weights (20)]} \qquad (21) \end{aligned}$$

- Left side: data. Right side: everything known from Step A, except D.
- Strictly decreasing in $D \Rightarrow$ invert for $D(y \mid s,x)$.
- Then $\delta(y \mid s,x) = \partial_y D(y \mid s,x)/\lambda_Y(y)$ [fundamental theorem of calculus].
- This is the paper's eq. (10), divided by its value at $y \downarrow s$.

Two selection problems, both handled by knowing G

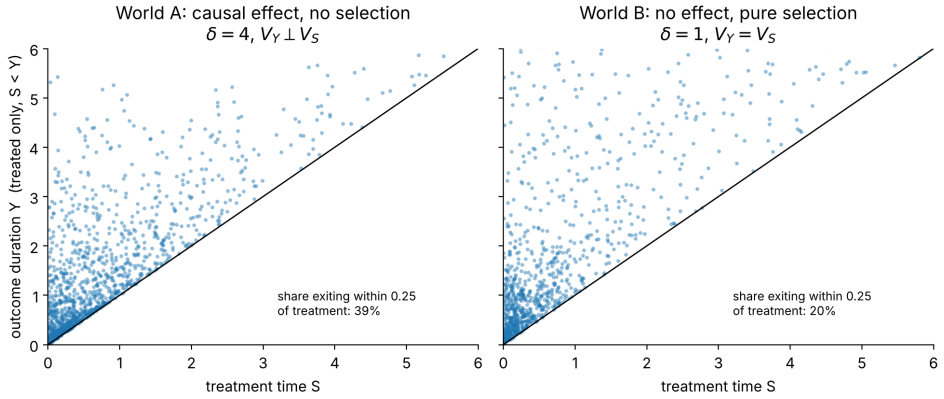
1. Selection *into* treatment at s: the v_S-tilt in (20).
2. Selection *out* after s: in (21), survivors of the treated cohort drift to low V_Y, as in Section 2.
   - Present whenever V_Y varies, even with random timing.
   - So naive hazard ratios at later y are not δ.

8.5 Summary: what is used where

→ *Why now:* collect which assumption did which job.

Assumption	Where it is used
No anticipation	Step A: before S there is no δ, so a clean competing-risks problem
Assumption 2	Steps A, B, D: independence given V; treated = risk set reweighted by v_S
MPH, $V \perp X$	Step A: everything is $\mathcal{L}$ of indices
$\mathbb{E}V < \infty$ (A4)	Step A, step 1: short-duration hazard ratios
Open set (A3)	Step A, step 2: trace $\mathcal{L}$ in 2-D $\Rightarrow$ dependence of V_Y, V_S

- Not used: an excluded instrument, selection on observables, parametric forms, multiple spells.



Intuition. Left: causal effect, no selection; exits bunch right after treatment at every S (local). Right: no effect, pure selection; S and Y correlated overall (global).

A Proof: any type distribution fits one survival curve

Setting of Section 2.5: data $\bar{S}$; true model $\bar{S}(y) = \mathcal{L}_G(\Lambda(y))$.

Construction: pick any $\tilde{G}$, then find the $\tilde{\Lambda}$ that fits. Take any $\tilde{G}$ with $\Pr(\tilde{V} = 0) = 0$. We need, for every y,

$$\mathcal{L}_{\tilde{G}}\big(\tilde{\Lambda}(y)\big) = \bar{S}(y). \qquad \text{[alternative model must reproduce the data]} \tag{22}$$

- At each y, (22) is **one equation in one unknown number** $\tilde{\Lambda}(y)$.
- $\mathcal{L}_{\tilde{G}}$ is continuous, strictly decreasing, from $\mathcal{L}(0) = 1$ to $\mathcal{L}(\infty) = \Pr(\tilde{V} = 0) = 0$ (Section 2.3).
  - So it takes every value in $(0, 1]$ exactly once, and (22) has a unique solution:

$$\begin{aligned}
&\tilde{\Lambda}(y) = \mathcal{L}_{\tilde{G}}^{-1}\big(\bar{S}(y)\big) && \text{[invert (22)]}\\
&\tilde{\Lambda}(0) = \mathcal{L}_{\tilde{G}}^{-1}(1) = 0 && [\bar{S}(0) = 1,\ \mathcal{L}(0) = 1]\\
&\quad \tilde{\Lambda} \text{ increasing, so } \tilde{\lambda} := \tilde{\Lambda}' \geq 0 && [\bar{S} \text{ and } \mathcal{L}_{\tilde{G}}^{-1} \text{ both decreasing}]
\end{aligned}$$

- $\tilde{\Lambda}$ is a valid integrated baseline, and $(\tilde{\Lambda}, \tilde{G})$ reproduces $\bar{S}$ exactly.
- This works for *every* $\tilde{G}$. So G, and hence the split, is not identified.

Mechanics in hazard terms. Differentiate (22):

$$\begin{aligned}
\bar{\theta}(y) &= -\frac{d}{dy}\log \mathcal{L}_{\tilde{G}}\big(\tilde{\Lambda}(y)\big) && \text{[(1); take } -\log \text{ of (22)]} \\
&= \tilde{\lambda}(y)\cdot\Big(-\frac{\mathcal{L}'_{\tilde{G}}}{\mathcal{L}_{\tilde{G}}}\Big)\big(\tilde{\Lambda}(y)\big) && \text{[chain rule; } \tilde{\Lambda}' = \tilde{\lambda}\text{]} \\
&= \tilde{\lambda}(y)\cdot\tilde{\mathbb{E}}[V \mid Y > y] && \text{[(9), computed under } \tilde{G}\text{]} \\
\implies \quad \tilde{\lambda}(y) &= \frac{\bar{\theta}(y)}{\tilde{\mathbb{E}}[V \mid Y > y]} && \text{[rearrange]}
\end{aligned}$$

- The observed $\bar{\theta}(y)$ is fixed. Each assumed $\tilde{G}$ implies some sorting path $\tilde{\mathbb{E}}[V \mid Y > y]$; $\tilde{\lambda}$ absorbs the rest.
- No heterogeneity ($\tilde{V} \equiv 1$): $\tilde{\mathbb{E}}[V \mid Y > y] = 1$, so $\tilde{\lambda} = \bar{\theta}$. The whole decline is duration dependence.
- More heterogeneity: $\tilde{\mathbb{E}}[V \mid Y > y]$ falls faster (Section 2.4: speed $\propto$ Var), so $\tilde{\lambda}$ has to fall less. More of the decline is attributed to sorting.
- With one duration, **trap vs. sorting is not identified**. The split is a pure assumption.